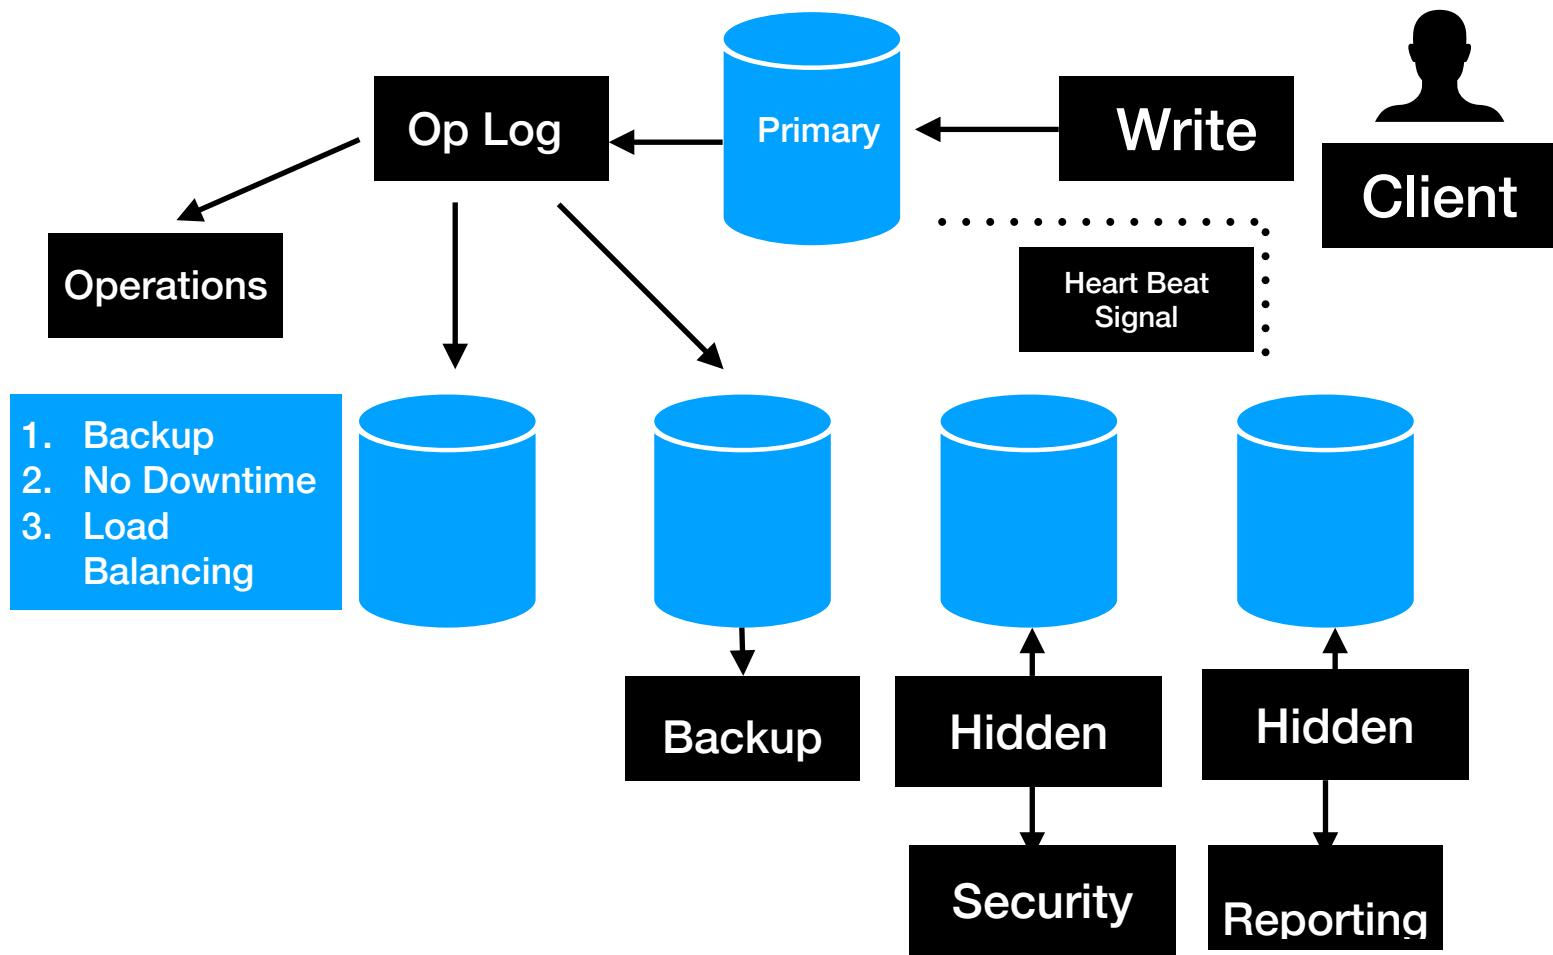


## Replica set:

1. Primary Instance can have Maximum of 50 Instance.
2. Primary & Secondary instance which is called Replica set.
3. Primary always send the Heart Beat Signal to Secondary and Vice Versa. Because of that both of know about their availability status.
4. All the Write Operation from the Client will always happen to the Primary Instance.
5. Whenever write operation happen Primary will generate an Op Log, and based one the configuration happens to the secondary Op Log will copy and execute. After execute this Op Log, secondary instance having the exact content of what the Primary having.
6. Actually, Once of the Secondary have reporting and Security. This is based on the Companies Interest.



## Setting Up Replica Set Servers:

Here, let assuming that, you have installed MongoDB Server and now you trying to create a db.

```
[centos@ip-172-26-8-241:~]$ mkdir replica01
[centos@ip-172-26-8-241:~]$ mkdir replica02
[centos@ip-172-26-8-241:~]$ mkdir replica03
[centos@ip-172-26-8-241:~]$ chmod -R 777 replica01
[centos@ip-172-26-8-241:~]$ chmod -R 777 replica02
[centos@ip-172-26-8-241:~]$ chmod -R 777 replica03
[centos@ip-172-26-8-241:~]$
```

```
[centos@ip-172-26-8-241:~]$ mongo --port 27020
MongoDB shell version v4.0.5
connecting to: mongodb://127.0.0.1:27020/?gssapiServiceName=mongodb
Implicit session: session { "id" : UUID("cf21f1bd-b4dd-4d1b-95b0-57c17f1898a9") }
MongoDB server version: 4.0.5
Server has startup warnings:
2019-02-04T04:33:11.046+0000 I CONTROL [initandlisten]
2019-02-04T04:33:11.046+0000 I CONTROL [initandlisten] ** WARNING: Access control is not enabled for the database
.
2019-02-04T04:33:11.046+0000 I CONTROL [initandlisten] ** Read and write access to data and configuration is unrestricted.
2019-02-04T04:33:11.046+0000 I CONTROL [initandlisten]
MongoDB Enterprise >
```

Now, I have created all three servers

```
[centos@ip-172-26-8-241:~]$ ps -ef | grep mongo
centos 16407 1 0 04:33 ? 00:00:01 mongod --dbpath replica01 --port 27020 --logpath replica01/log --logappend --fork --bind_ip_all --replSet myreplicaset01
centos 16460 1 1 04:35 ? 00:00:01 mongod --dbpath replica02 --port 27021 --logpath replica02/log --logappend --fork --bind_ip_all --replSet myreplicaset01
centos 16491 1 6 04:36 ? 00:00:00 mongod --dbpath replica03 --port 27022 --logpath replica03/log --logappend --fork --bind_ip_all --replSet myreplicaset01
centos 16520 16375 0 04:36 pts/0 00:00:00 grep --color=auto mongo
[centos@ip-172-26-8-241:~]$
```

Initiating Replica set and adding members.

```
MongoDB Enterprise > rs.initiate()
{
    "info2" : "no configuration specified. Using a default configuration for the set",
    "me" : "ip-172-26-8-241.ap-south-1.compute.internal:27020",
    "ok" : 1,
    "operationTime" : Timestamp(1549255229, 1),
    "$clusterTime" : {
        "clusterTime" : Timestamp(1549255229, 1),
        "signature" : {
            "hash" : BinData(0,"AAAAAAAAAAAAAAAAAAAAA="),
            "keyId" : NumberLong(0)
        }
    }
}
MongoDB Enterprise myreplicaset01:SECONDARY>
MongoDB Enterprise myreplicaset01:PRIMARY>
```

When you fire this command, rs.config() -> you can see at this point of time only one member is there. And, also you can rs.status().

Now, I am adding members. After successful adding we can check the rs.config() and rs.status()

```
MongoDB Enterprise myreplicaset01:PRIMARY> rs.add("ip-172-26-8-241.ap-south-1.compute.internal:27021")
{
    "ok" : 1,
    "operationTime" : Timestamp(1549255355, 1),
    "$clusterTime" : {
        "clusterTime" : Timestamp(1549255355, 1),
        "signature" : {
            "hash" : BinData(0, "AAAAAAAAAAAAAAAAAAAAA="),
            "keyId" : NumberLong(0)
        }
    }
}

MongoDB Enterprise myreplicaset01:PRIMARY> rs.add("ip-172-26-8-241.ap-south-1.compute.internal:27022")
{
    "ok" : 1,
    "operationTime" : Timestamp(1549255414, 1),
    "$clusterTime" : {
        "clusterTime" : Timestamp(1549255414, 1),
        "signature" : {
            "hash" : BinData(0, "AAAAAAAAAAAAAAAAAAAAA="),
            "keyId" : NumberLong(0)
        }
    }
}
MongoDB Enterprise myreplicaset01:PRIMARY>
```

Now, I am adding at add documents in Primary.

```
[centos@ip-172-26-8-241:~]
MongoDB Enterprise myreplicaset01:PRIMARY> use employees
switched to db employees
MongoDB Enterprise myreplicaset01:PRIMARY> db.myemployees.insert({_id:1, name:"myemp1", sal:500});
WriteResult({ "nInserted" : 1 })
MongoDB Enterprise myreplicaset01:PRIMARY> db.myemployees.find();
{ "_id" : 1, "name" : "myemp1", "sal" : 500 }
MongoDB Enterprise myreplicaset01:PRIMARY>
```

After successfully added, Now, I am going to check from secondary server.

```
[centos@ip-172-26-8-241 ~]$ mongo --port 27021
MongoDB shell version v4.0.5
connecting to: mongodb://127.0.0.1:27021/?gssapiServiceName=mongodb
Implicit session: session { "id" : UUID("6bc017e8-3364-4848-a30f-b7114b2e7f8e") }
MongoDB server version: 4.0.5
Server has startup warnings:
2019-02-04T04:35:10.642+0000 I CONTROL  [initandlisten]
2019-02-04T04:35:10.643+0000 I CONTROL  [initandlisten] ** WARNING: Access control is not enabled for the database.
.
2019-02-04T04:35:10.643+0000 I CONTROL  [initandlisten] ** Read and write access to data and configuration is unrestricted.
2019-02-04T04:35:10.643+0000 I CONTROL  [initandlisten]
MongoDB Enterprise myreplicaset01:SECONDARY> show dbs;
2019-02-04T04:50:01.437+0000 E QUERY    [js] Error: listDatabases failed:
    "operationTime" : Timestamp(1549255801, 1),
    "ok" : 0,
    "errmsg" : "not master and slaveOk=false",
    "code" : 13435,
    "codeName" : "NotMasterNoSlaveOk",
    "$clusterTime" : {
        "clusterTime" : Timestamp(1549255801, 1),
        "signature" : {
            "hash" : BinData(0, "AAAAAAAAAAAAAAAAAAAAA="),
            "keyId" : NumberLong(0)
        }
    }
} :
_getErrorWithCode@src/mongo/shell/utils.js:25:13
Mongo.prototype.getDBs@src/mongo/shell/mongo.js:124:1
shellHelper.show@src/mongo/shell/utils.js:876:19
shellHelper@src/mongo/shell/utils.js:766:15
```

I am getting this error because by default you will not access Secondary node.

Please Note: Write Operation Always Primary and Read Operation by default is Primary. You need to set slaveOk = true Changs in the Secondary server then, you can actually go ahead and do the read operations

```
centos@ip-172-26-8-241:~  
MongoDB Enterprise myreplicaset01:SECONDARY> rs.slaveOk(true);  
MongoDB Enterprise myreplicaset01:SECONDARY> show dbs;  
admin      0.000GB  
config     0.000GB  
employees  0.000GB  
local      0.000GB  
MongoDB Enterprise myreplicaset01:SECONDARY> █
```

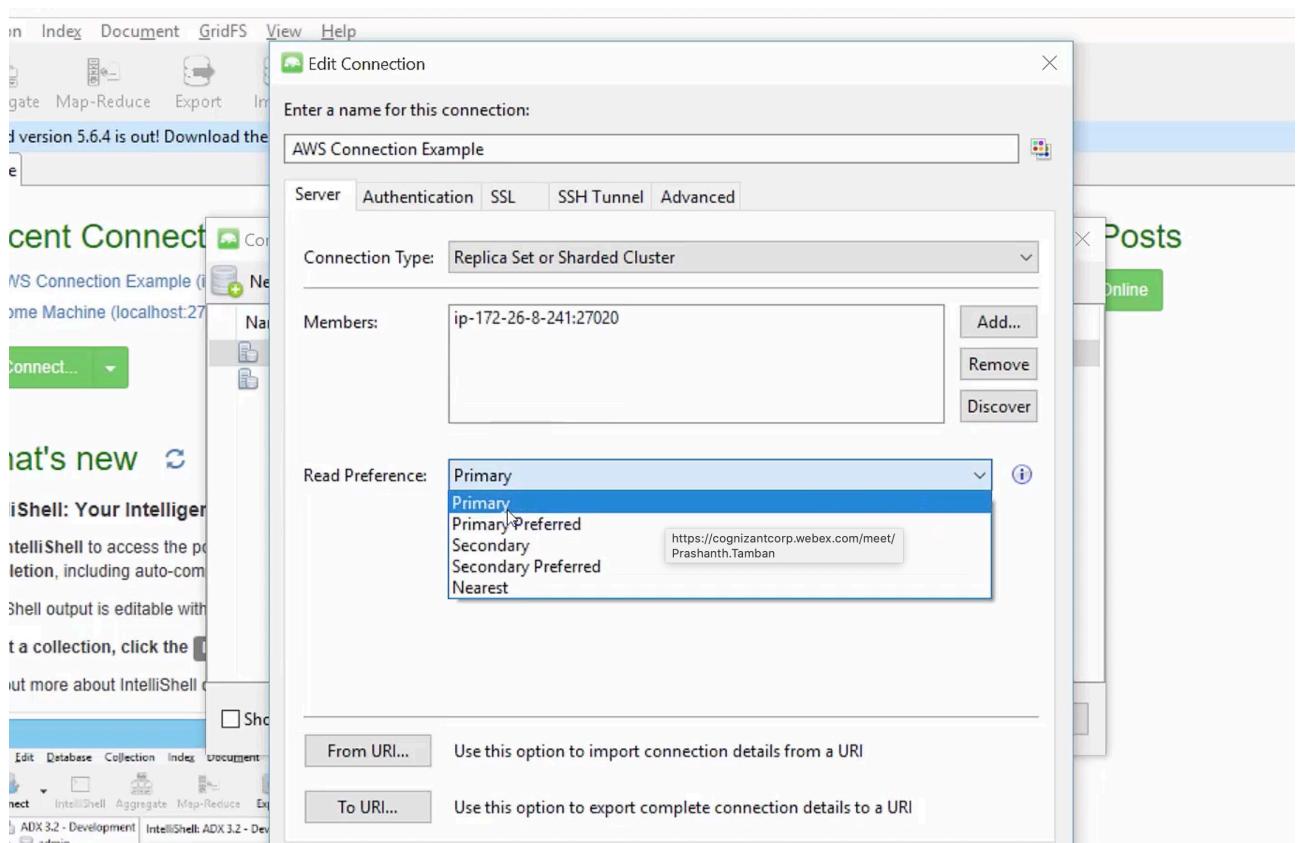
Now, you see show dbs; shows the output and find the myemployees output.

Importantly, Secondary Server Only Read Operation is allowed. And you must need to set rs.slaveOk(true) for every secondary node.

```
MongoDB Enterprise myreplicaset01:SECONDARY> rs.slaveOk(true);  
MongoDB Enterprise myreplicaset01:SECONDARY> show dbs;  
admin      0.000GB  
config     0.000GB  
employees  0.000GB  
local      0.000GB  
MongoDB Enterprise myreplicaset01:SECONDARY> use employees;  
switched to db employees  
MongoDB Enterprise myreplicaset01:SECONDARY> db.myemployees.find()  
{ "_id" : 1, "name" : "myempl", "sal" : 500 }  
MongoDB Enterprise myreplicaset01:SECONDARY> █
```

now, I am connecting through MongoChef Core - 3T, which IDE for MongoDB. Here, you have Read Preference. Meaning, first you have to set rs.slaveOk(true) and later when read operation comes, you check actually select the Read Preference which close to Primary.

Hence, I request you to do the same as above in your MongoDB and check the output and after that do the necessary changes.



MongoChef Core - 3T Software Labs

File Edit Database Collection Index Document GridFS View Help

Connect IntelliShell Aggregate Map-Reduce Export

MongoChef is now [Studio 3T](#), and version 5.6.4 is out! Download the update [here](#) to experience the latest features!

**AWS Connection Example [replica set]**

- Replica Set Members
  - ip-172-26-8-241.ap-south-1.co
  - ip-172-26-8-241.ap-south-1.co
  - ip-172-26-8-241.ap-south-1.co
- admin
- config
- employees**
  - Database: employees
  - Collections: 1
  - Objects: 1
  - Data Size: 48 B (48)
  - Storage Size: 16.0 KiB (16,384)
  - Indexes: 1
  - Index Size: 16.0 KiB (16,384)
- local

Welcome

## Recent Connections

- AWS Connection Example ((3 servers))
- Home Machine (localhost:27017)

**Connect...**

## What's new

**IntelliShell: Your Intelligent MongoDB Shell**

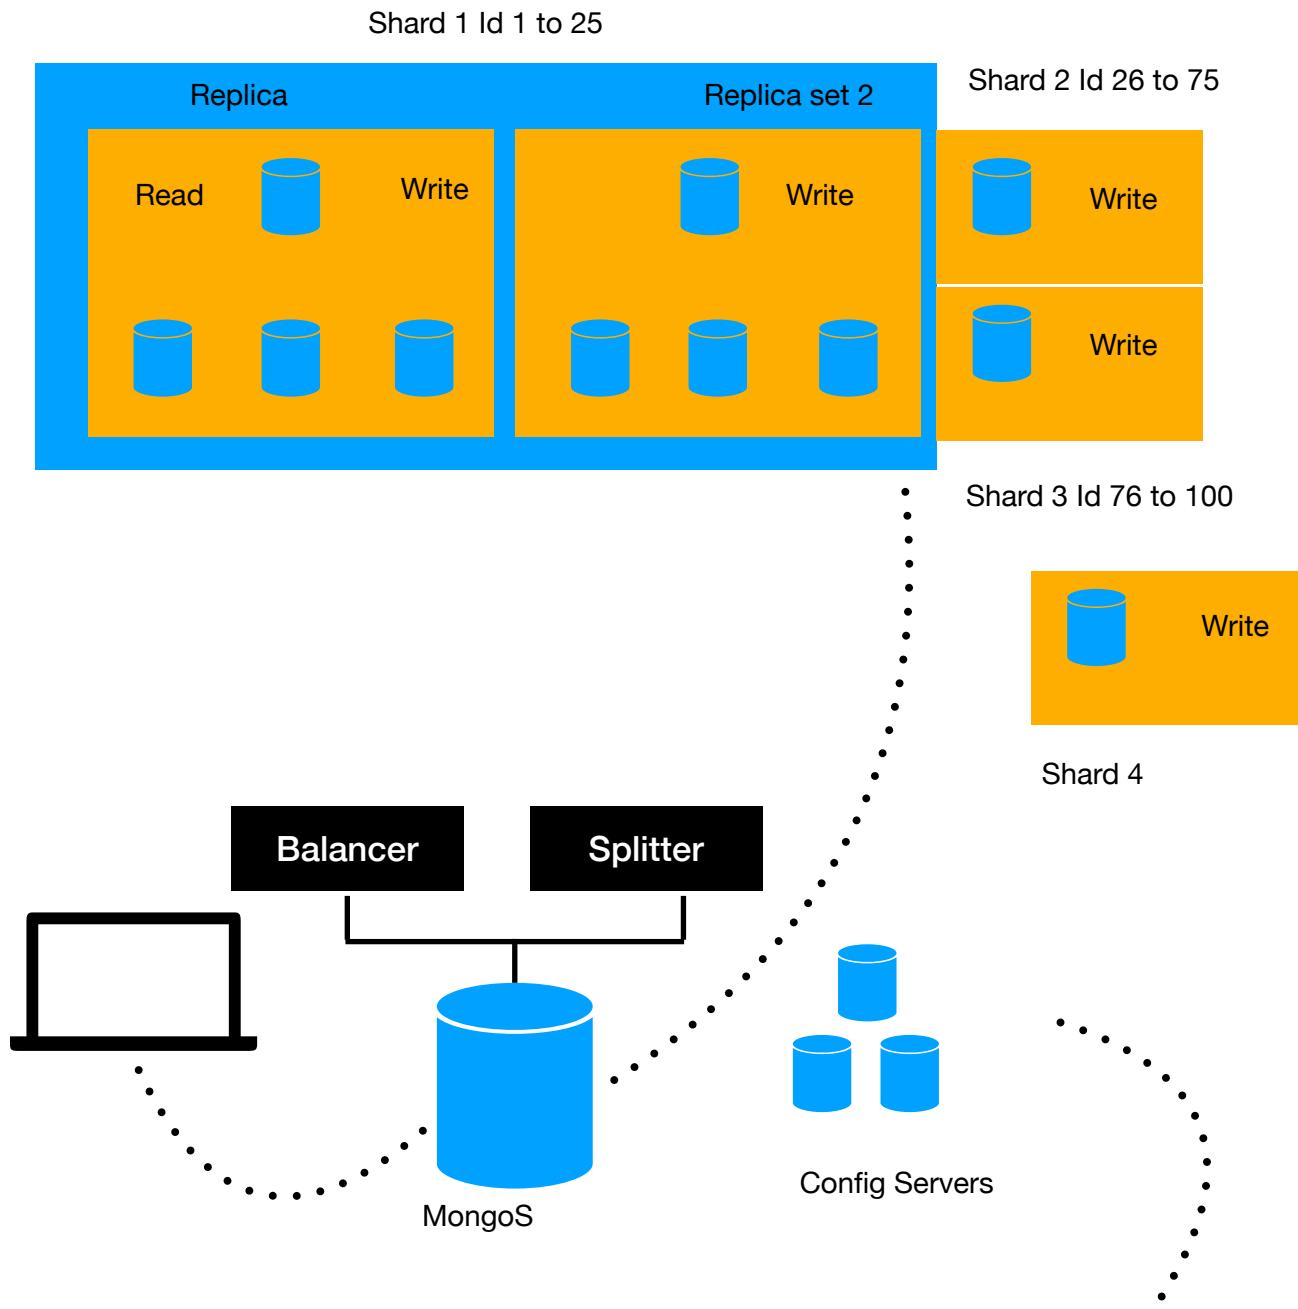
Use **IntelliShell** to access the power of the shell in the beautiful MongoChef environment. Enjoy intelligent auto-completion, including auto-completion of shell methods and document field names.

IntelliShell output is editable with MongoChef's in-place editors and can always be exported to a collection.

Select a collection, click the **IntelliShell** button in the main tool bar and check it out!

Find out more about **IntelliShell** on our blog: [Getting Started with MongoChef's IntelliShell](#)

## Sharding Techniques:



Shard 1	1	25
Shard 2	26	75
Shard 3	76	100

1. Shard 1 has data from 1 to 25, Shard 2 has data from 26 to 75, Shard 3 has data from 76 to 100. This information is mentioned in Config Servers.
2. Whenever any request from Customer, it will go to MongoS and contacting the Config Server and get the location based on which the request has been routed to either Shard 1 or 2 or 3.
3. Balancer and Splitter will ensure that, at any given point of time all the Shards are equally utilised.
4. If you are introducing the new Shard 4, then you don't have to worry because Balancer and Splitter will find that Shard 4 is newly added and which under performed. So then, all the data will be shared to Shard 4 and make them ready to serve the request. Which is called automatic scaling.
5. If you are simply calling `find()` then, mongos doesn't know which Shard needs to call, therefore Mongos broadcasts it to all Shards and finds it. This is also possible to reduce the performance.
6. Always, Operation contains Shard Key, request goes to the respected Shard and fetches the data.

**Important:** Whenever your operations are Read Intensive, you can have Replica set at your rescue. Whenever your operations are Write intensive, you can have Shard Cluster at your rescue. Together you have Read & Write powerful mechanism

#### Balancer & Splitter:

1. For e.g. Shard 1 has 10 GB Memory, Shard 2 & 3 also having 10 GB Memory. Assuming here, Shard 1 receiving multiple requests and Shard 1's memory consumption is 8 GB. In the given situation, Balancer & Splitter understand that other 2 Shards are not utilised hence, Shard 1's data will be split to two Shards so that, utilisation is equal at any given point of time.

#### Capped Collections:

When you create a collection, if you set `capped = true`, then, when retrieval process, data will be retrieved the way you intended. Non-capped collection, there is no guarantee that retrieval may require you to use `Natural`.